

# **CIS 655 Advanced Computer Architecture**

## **Fall Final Report**

**Group Member**

**Name (SUID)**

**LIU, SHOU YUNG (514096323)**

**Zezhi Su(368202764)**

1. Project Scope

In the final project, we modify the project scope to match the project requirement. Below are the target which we focus on.

- 1. Write a module to intercept all the page faults and send signal to the user process.
- 2. Find the mechanism to trigger original page fault function

2. Project member

SUID	Name	Email
368202764	Zezhi Su	zsu103@syr.edu
514096323	LIU, Shou Yung	sliu132@syr.edu

3. Project Description

1. Introduction

Following the project requirements, we need to find the mechanism to intercept the page fault occur, and also trigger original page fault mechanism. The difficult point of the project is “How to invoke the original do page fault function without modify the original kernel module?” If we could modify the original kernel function, then it will be easier to achieve what we expect to. Modifying the do page fault function and declare the function to extern, then write another Linux Kernel Module to call the original page fault function after print the page fault address detail information. But if we want to get the page fault information also trigger original do page fault, we have to find another way to achieve this. Our solution is reading the function from /proc/kallsyms, which is the Linux Kernel global symbol table. Extract the non-export function address from the symbol table, so that we could execute the original page fault mechanism. Next section describe the purpose of the file. The third section display the execution screen shot and code explanation.

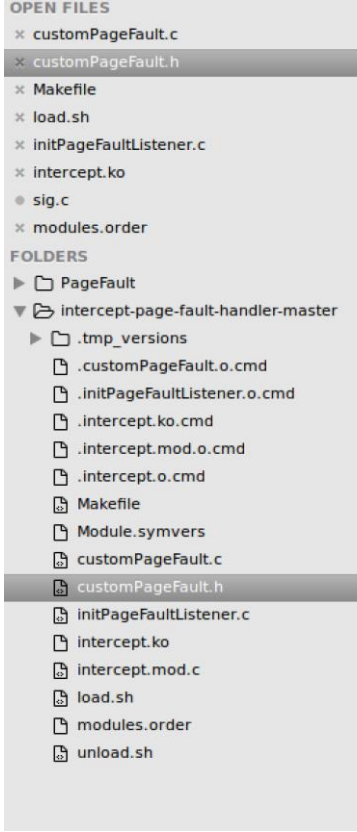
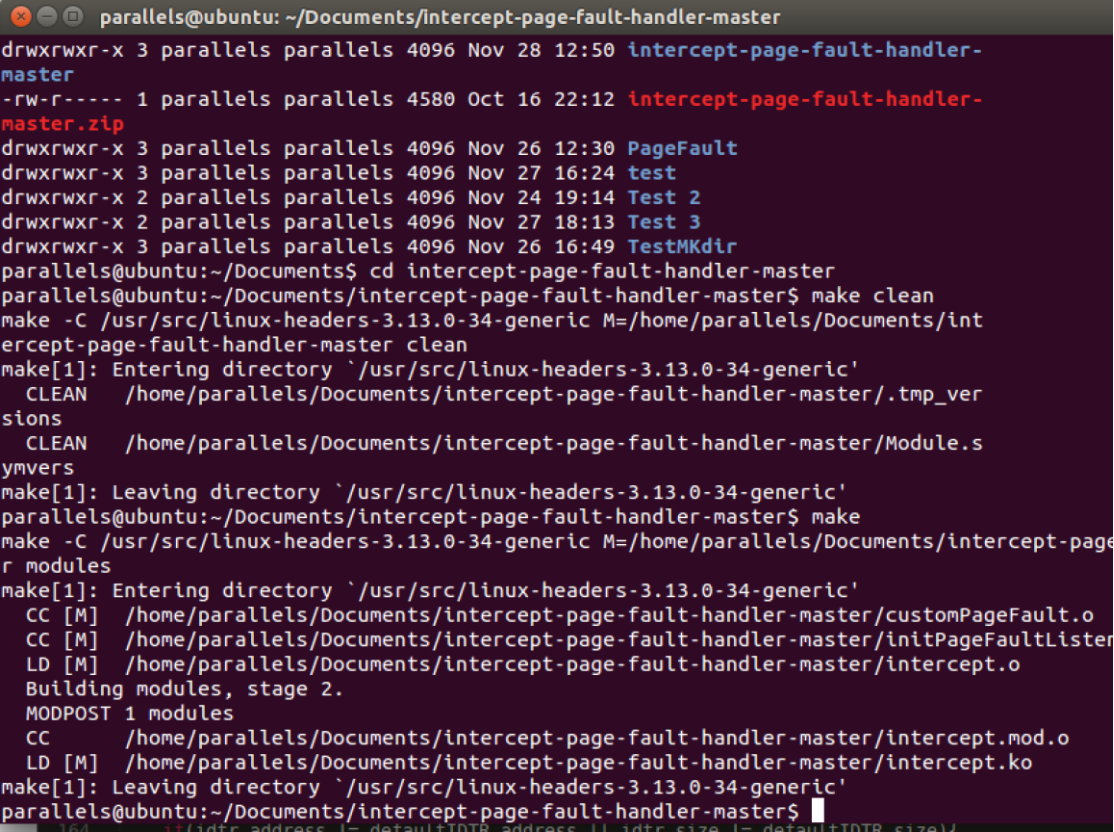
2. Project file Description

File List	Description
load.sh	<ul style="list-style-type: none"><li>1. Load the kernel function from the <b>boot/System.map-\${LINUX_VERSION}</b>, which could correctly execute the whole project.</li><li>2. Linked compiled module to the kernel</li></ul>
unload.sh	Remove linked module from the kernel
makefile	<ul style="list-style-type: none"><li>1. Compile customPageFault.c, initPageFaultListener.c two file into intercept.o, also generate intercept.ko, intercept.mod.c, module.order, files which are needed for the kernel module.</li><li>2. Clean compiled module</li></ul>
initPageFaultListener.c	<p>Include two function- init_page_fault_listener, exit_page_fault_listener</p> <ul style="list-style-type: none"><li>1. init_page_fault_listener: register customPageFault function.</li><li>2. exit_page_fault_listener: remove kernel module</li></ul>
customPageFault.c	<ul style="list-style-type: none"><li>1. registerPageFaultListener:<ul style="list-style-type: none"><li>• Register page fault function. Receive parameters which comes from symbol table.</li><li>• Store default idtr</li><li>• Get the idtr default address</li><li>• Allocate new page for new idtr</li><li>• Copy the old idtr to new page</li></ul></li></ul>

	<ul style="list-style-type: none"> <li>Loaded idt to all the cpu</li> </ul> 2. unregisterPageFaultListener: <ul style="list-style-type: none"> <li>restore default idtr and free the allocate page</li> </ul> 3. write_pid <ul style="list-style-type: none"> <li>Get user space process id from then send message to it</li> </ul>
customPageFault.h	customPageFault.h Header file
UserSpaceCode.c	UserSpace process which could receive the kernel information.

### 3. Execution Step

- Under Ubuntu 3.13.0-34-generic X86\_64 platform open terminal
- Change directory to the project directory
- Type “make”

Make result	Step from 1 to 3
	

- Type `sudo ./load.sh`
- Successful insert module from it

```

parallels@ubuntu: ~/Documents/intercept-page-fault-handler-master
drwxrwxr-x 3 parallels parallels 4096 Nov 26 12:30 PageFault
drwxrwxr-x 3 parallels parallels 4096 Nov 27 16:24 test
drwxrwxr-x 2 parallels parallels 4096 Nov 24 19:14 Test 2
drwxrwxr-x 2 parallels parallels 4096 Nov 27 18:13 Test 3
drwxrwxr-x 3 parallels parallels 4096 Nov 26 16:49 TestMKDir
parallels@ubuntu:~/Documents$ cd intercept-page-fault-handler-master
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ make clean
make -C /usr/src/linux-headers-3.13.0-34-generic M=/home/parallels/Documents/int
ercept-page-fault-handler-master clean
make[1]: Entering directory `/usr/src/linux-headers-3.13.0-34-generic'
CLEAN /home/parallels/Documents/intercept-page-fault-handler-master/.tmp_ver
sions
CLEAN /home/parallels/Documents/intercept-page-fault-handler-master/Module.s
ymvers
make[1]: Leaving directory `/usr/src/linux-headers-3.13.0-34-generic'
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ make
make -C /usr/src/linux-headers-3.13.0-34-generic M=/home/parallels/Documents/intercept-page-fault-handler-maste
r modules
make[1]: Entering directory `/usr/src/linux-headers-3.13.0-34-generic'
CC [M] /home/parallels/Documents/intercept-page-fault-handler-master/customPageFault.o
CC [M] /home/parallels/Documents/intercept-page-fault-handler-master/initPageFaultListener.o
LD [M] /home/parallels/Documents/intercept-page-fault-handler-master/intercept.o
Building modules, stage 2.
MODPOST 1 modules
CC /home/parallels/Documents/intercept-page-fault-handler-master/intercept.mod.o
LD [M] /home/parallels/Documents/intercept-page-fault-handler-master/intercept.ko
make[1]: Leaving directory `/usr/src/linux-headers-3.13.0-34-generic'
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ sudo ./load.sh
[sudo] password for parallels:
Sorry, try again.
[sudo] password for parallels:
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ █

```

- vi. Type “dmesg | tail” to show the printk page fault information

```

parallels@ubuntu: ~/Documents/intercept-page-fault-handler-master
sions
CLEAN /home/parallels/Documents/intercept-page-fault-handler-master/Module.s
ymvers
make[1]: Leaving directory `/usr/src/linux-headers-3.13.0-34-generic'
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ make
make -C /usr/src/linux-headers-3.13.0-34-generic M=/home/parallels/Documents/intercept-page-fault-handler-maste
r modules
make[1]: Entering directory `/usr/src/linux-headers-3.13.0-34-generic'
CC [M] /home/parallels/Documents/intercept-page-fault-handler-master/customPageFault.o
CC [M] /home/parallels/Documents/intercept-page-fault-handler-master/initPageFaultListener.o
LD [M] /home/parallels/Documents/intercept-page-fault-handler-master/intercept.o
Building modules, stage 2.
MODPOST 1 modules
CC /home/parallels/Documents/intercept-page-fault-handler-master/intercept.mod.o
LD [M] /home/parallels/Documents/intercept-page-fault-handler-master/intercept.ko
make[1]: Leaving directory `/usr/src/linux-headers-3.13.0-34-generic'
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ sudo ./load.sh
[sudo] password for parallels:
Sorry, try again.
[sudo] password for parallels:
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ dmesg | tail
[ 8443.668369] Page fault Listener : page fault detected in process 29286.
[ 8443.668372] Page fault Listener : page fault detected in process 29286.
[ 8443.668380] Page fault Listener : page fault detected in process 29286.
[ 8443.668383] Page fault Listener : page fault detected in process 29286.
[ 8443.668387] Page fault Listener : page fault detected in process 29286.
[ 8443.668391] Page fault Listener : page fault detected in process 29286.
[ 8443.668394] Page fault Listener : page fault detected in process 29286.
[ 8443.668397] Page fault Listener : page fault detected in process 29286.
[ 8443.668399] Page fault Listener : page fault detected in process 29286.
[ 8443.668406] Page fault Listener : page fault detected in process 29286.
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$

```

- vii. Type ./unload.sh to remove module, and use “lsmod” check the module is indeed remove



```

parallels@ubuntu: ~/Documents/intercept-page-fault-handler-master
make -C /usr/src/linux-headers-3.13.0-34-generic M=/home/parallels/Documents/intercept-page-fault-handler-master modules
make[1]: Entering directory `/usr/src/linux-headers-3.13.0-34-generic'
CC [M] /home/parallels/Documents/intercept-page-fault-handler-master/customPageFault.o
CC [M] /home/parallels/Documents/intercept-page-fault-handler-master/initPageFaultListener.o
LD [M] /home/parallels/Documents/intercept-page-fault-handler-master/intercept.o
Building modules, stage 2.
MODPOST 1 modules
CC /home/parallels/Documents/intercept-page-fault-handler-master/intercept.mod.o
LD [M] /home/parallels/Documents/intercept-page-fault-handler-master/intercept.ko
make[1]: Leaving directory `/usr/src/linux-headers-3.13.0-34-generic'
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ sudo ./load.sh
[sudo] password for parallels:
Sorry, try again.
[sudo] password for parallels:
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ dmesg | tail
[ 8443.668369] Page fault Listener : page fault detected in process 29286.
[ 8443.668372] Page fault Listener : page fault detected in process 29286.
[ 8443.668380] Page fault Listener : page fault detected in process 29286.
[ 8443.668383] Page fault Listener : page fault detected in process 29286.
[ 8443.668387] Page fault Listener : page fault detected in process 29286.
[ 8443.668391] Page fault Listener : page fault detected in process 29286.
[ 8443.668394] Page fault Listener : page fault detected in process 29286.
[ 8443.668397] Page fault Listener : page fault detected in process 29286.
[ 8443.668399] Page fault Listener : page fault detected in process 29286.
[ 8443.668406] Page fault Listener : page fault detected in process 29286.
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ ./unload.sh
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ lsmod
Module                Size  Used by
nls_utf8              12557  1
isofs                 39835  1
usb_lip               22891  0

```

viii. Execute UserSpaceCode to receive the signal send from kernel space

```

parallels@ubuntu: ~/Documents/intercept-page-fault-handler-master
dm_multipath          22873  0
cryptd                20359  3 ghash_c1mulni_intel,aesni_intel,ablk_helper
uvcvideo              80885  0
snd_timer             29482  2 snd_pcm,snd_seq
videobuf2_vmalloc     13216  1 uvcvideo
scsi_dh               14882  1 dm_multipath
videobuf2_memops      13362  1 videobuf2_vmalloc
snd                   69238  12 snd_ac97_codec,snd_intel8x0,snd_timer,snd_pcm,s
nd_seq,snd_rawmidi,snd_seq_device,snd_seq_midi
videobuf2_core        40664  1 uvcvideo
videodev              134688  2 uvcvideo,videobuf2_core
serio_raw             13462  0
soundcore             12680  1 snd
rfcomm                69160  0
bnep                  19624  2
bluetooth             391196  10 bnep,rfcomm
lpc_ich               21080  0
pvpanic               12801  0
shpchp                37032  0
prl_tg                21944  1 prl_fs
parport_pc            32701  0
ppdev                17671  0
lp                    17759  0
parport               42348  3 lp,ppdev,parport_pc
mac_hid               13205  0
psmouse               106678  0
ahci                  25819  3
libahci               32560  1 ahci
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ sudo ./a.out
received value 1234
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$

```

ix. Show that intercept kernel still in kernel module

```
parallels@ubuntu: ~/Documents/Intercept-page-fault-handler-master
snd_timer                29482    2  snd_pcm,snd_seq
videobuf2_vmalloc        13216    1  uvcvideo
scsi_dh                   14882    1  dm_multipath
videobuf2_memops         13362    1  videobuf2_vmalloc
snd                       69238   12  snd_ac97_codec,snd_intel8x0,snd_timer,snd_pcm,s
nd_seq,snd_rawmidi,snd_seq_device,snd_seq_midi
videobuf2_core           40664    1  uvcvideo
videodev                 134688    2  uvcvideo,videobuf2_core
serio_raw                13462    0
soundcore                12680    1  snd
rfcomm                   69160    0
bnep                     19624    2
bluetooth                391196   10  bnep,rfcomm
lpc_ich                  21080    0
pvpanic                  12801    0
shpchp                   37032    0
prl_tg                   21944    1  prl_fs
parport_pc               32701    0
ppdev                    17671    0
lp                       17759    0
parport                  42348    3  lp,ppdev,parport_pc
mac_hid                  13205    0
psmouse                  106678   0
ahci                     25819    3
libahci                  32560    1  ahci
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ sudo ./a.out
received value 1234
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$ lsmod |grep in
tercept
intercept                13690    0
parallels@ubuntu:~/Documents/intercept-page-fault-handler-master$
```

Show that user process receive the kernel signal

#### 4. Project Critical Point Explanation

There are some point worth to mention in this project. First of all, figuring out the Page Fault process. How the process begin, how the stack information store to handle the page fault exception. Then find the register which store the address we need. For example, `rsp`, represent stack segment in 64 bit architecture. “`rdi`” register represent index and pointer in 64 architecture. Secondly, figure out the page fault process and every detail in different variables, such as error code, copy-on-write characteristic in idt table, how to get function address from the global symbol table. Third, how to use assembly language to change the default page fault behavior. For our team members, it is a challenge task to write some assembly code. Since we have to very carefully deal with copy register value, also, without effect the system operation. Finally, how to send process to the specific user process. We use the “`debugfs`” to get the user process id, then send the message from our page fault handler. In the future, we expect to do some modification when the kernel send the signal to process, then resume to wait for user process response.

#### 5. Conclusion

Intercept page fault in Linux kernel is not an easy task to accomplish with lack of knowledge in kernel operation. So choosing this topic is quite challenge to our team members. But in the whole project, we spend most of time to figure out the process in Linux Page fault mechanism, signal handle, and how the whole process is ongoing when the page fault exception occurs. We did gain a lot precious experience from the project. The more time we spend to research in kernel module. The more familiar with the computer architecture. Expecting the next semester, we could go deeper in the Linux kernel programming.

## 6. Reference

- i. Daniel P Bovet, Marco Cesati. (2005). *Understanding the Linux kernel*, 3rd edition. O'Reily.
- ii. Keller, A. (2008). *Kernel Space - User Space Interfaces*. Retrieved from [http://people.ee.ethz.ch/~arkeller/linux/multi/kernel\\_user\\_space\\_howto-5.html](http://people.ee.ethz.ch/~arkeller/linux/multi/kernel_user_space_howto-5.html)
- iii. LXR. (n.d.). *Linux Cross Reference*. Retrieved from [http://lxr.free-electrons.com/source/drivers/infiniband/ulp/iser/iscsi\\_iser.c#L101](http://lxr.free-electrons.com/source/drivers/infiniband/ulp/iser/iscsi_iser.c#L101).
- iv. Molly, D. (2015, 4 14). *Part 1: Introduction*. Retrieved from Writing a Linux kernel Module: <http://derekmolloy.ie/writing-a-linux-kernel-module-part-1-introduction/>
- v. Peter J Salzman, Michael Burian, and Ori Pomerantz. (n.d.). *Writing Your Own Loadable Kernel Module*. Retrieved from <http://www.tldp.org/>: <http://www.tldp.org/HOWTO/Module-HOWTO/x839.html>
- vi. Phrack, K. (2003, 8 13). *Handling Interrupt Descriptor Table for Fun and Profit*. Retrieved from Handling Interrupt Descriptor Table for Fun and Profit: <http://www.phrack.org>
- vii. Pommnitz, J. (n.d.). *Kernel level exception handling in Linux Commentary*. Retrieved from <https://www.kernel.org/>: <http://www.kernel.org/doc/Documentation/x86/exception-tables.txt>
- viii. USTC, R. (2013, 04 25). *CNBlog*. Retrieved from <http://www.cnblogs.com/richardustc/archive/2013/04/25/3043674.html>
- ix. VBird. (n.d.). *VBird WebSite*. Retrieved from V Bird: [http://linux.vbird.org/linux\\_basic/redhat6.1/linux\\_06command.php#grep](http://linux.vbird.org/linux_basic/redhat6.1/linux_06command.php#grep)
- x. ZHANG, S. (2011, 03 04). *May the #! with you*. Retrieved from <https://onebitbug.me/2011/03/04/introducing-linux-kernel-symbols/>.